

```
In [53]: import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
from scipy import stats
import statsmodels.api as sm
import warnings
warnings.filterwarnings("ignore")

# Loading the dataset
df_thyroid = pd.read_csv('/Users/roshan/Downloads/thyroid_cancer_risk_data.c

# Print some data
print(df_thyroid.info())
print(df_thyroid.head())

# Check for missing values
print(df_thyroid.isnull().sum())

# Summary statistics
print(df_thyroid.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 212691 entries, 0 to 212690
```

```
Data columns (total 17 columns):
```

#	Column	Non-Null Count	Dtype
0	Patient_ID	212691 non-null	int64
1	Age	212691 non-null	int64
2	Gender	212691 non-null	object
3	Country	212691 non-null	object
4	Ethnicity	212691 non-null	object
5	Family_History	212691 non-null	object
6	Radiation_Exposure	212691 non-null	object
7	Iodine_Deficiency	212691 non-null	object
8	Smoking	212691 non-null	object
9	Obesity	212691 non-null	object
10	Diabetes	212691 non-null	object
11	TSH_Level	212691 non-null	float64
12	T3_Level	212691 non-null	float64
13	T4_Level	212691 non-null	float64
14	Nodule_Size	212691 non-null	float64
15	Thyroid_Cancer_Risk	212691 non-null	object
16	Diagnosis	212691 non-null	object

```
dtypes: float64(4), int64(2), object(11)
```

```
memory usage: 27.6+ MB
```

```
None
```

	Patient_ID	Age	Gender	Country	Ethnicity	Family_History	\
0	1	66	Male	Russia	Caucasian	No	
1	2	29	Male	Germany	Hispanic	No	
2	3	86	Male	Nigeria	Caucasian	No	
3	4	75	Female	India	Asian	No	
4	5	35	Female	Germany	African	Yes	

	Radiation_Exposure	Iodine_Deficiency	Smoking	Obesity	Diabetes	TSH_Level
0	Yes	No	No	No	No	9.37
1	Yes	No	No	No	No	1.83
2	No	No	No	No	No	6.26
3	No	No	No	No	No	4.10
4	Yes	No	No	No	No	9.10

	T3_Level	T4_Level	Nodule_Size	Thyroid_Cancer_Risk	Diagnosis
0	1.67	6.16	1.08	Low	Benign
1	1.73	10.54	4.05	Low	Benign
2	2.59	10.57	4.61	Low	Benign
3	2.62	11.04	2.46	Medium	Benign
4	2.11	10.71	2.11	High	Benign

Patient_ID	0
Age	0
Gender	0
Country	0
Ethnicity	0
Family_History	0
Radiation_Exposure	0
Iodine_Deficiency	0
Smoking	0
Obesity	0

```

Diabetes      0
TSH_Level     0
T3_Level     0
T4_Level     0
Nodule_Size   0
Thyroid_Cancer_Risk  0
Diagnosis     0
dtype: int64

```

	Patient_ID	Age	TSH_Level	T3_Level \
count	212691.00000	212691.000000	212691.000000	212691.000000
mean	106346.00000	51.918497	5.045102	2.001727
std	61398.74739	21.632815	2.860264	0.866248
min	1.00000	15.000000	0.100000	0.500000
25%	53173.50000	33.000000	2.570000	1.250000
50%	106346.00000	52.000000	5.040000	2.000000
75%	159518.50000	71.000000	7.520000	2.750000
max	212691.00000	89.000000	10.000000	3.500000

	T4_Level	Nodule_Size
count	212691.000000	212691.000000
mean	8.246204	2.503403
std	2.164188	1.444631
min	4.500000	0.000000
25%	6.370000	1.250000
50%	8.240000	2.510000
75%	10.120000	3.760000
max	12.000000	5.000000

```

In [13]: # Select relevant variables
selected_vars = ["Age", "Nodule_Size", "Radiation_Exposure", "Family_History"]

# Encode categorical variables
df_thyroid_encoded = df_thyroid.copy()
df_thyroid_encoded["Radiation_Exposure"] = df_thyroid_encoded["Radiation_Exposure"].map({"Benign": 0, "Malignant": 1})
df_thyroid_encoded["Family_History"] = df_thyroid_encoded["Family_History"].map({"No": 0, "Yes": 1})
df_thyroid_encoded["Diagnosis"] = df_thyroid_encoded["Diagnosis"].map({"Benign": 0, "Malignant": 1})

# Compute descriptive statistics
desc_stats = df_thyroid_encoded[selected_vars].describe()

# Compute mode for each variable
mode_values = df_thyroid_encoded[selected_vars].mode().iloc[0]

# Append mode row to descriptive statistics
desc_stats.loc["Mode"] = mode_values
# Print descriptive statistics
print("### Descriptive Statistics: Mean, Mode, Spread, and Tails ###")
print(desc_stats)

```

Descriptive Statistics: Mean, Mode, Spread, and Tails

	Age	Nodule_Size	Radiation_Exposure	Family_History \
count	212691.000000	212691.000000	212691.000000	212691.000000
mean	51.918497	2.503403	0.149795	0.300083
std	21.632815	1.444631	0.356871	0.458295
min	15.000000	0.000000	0.000000	0.000000
25%	33.000000	1.250000	0.000000	0.000000
50%	52.000000	2.510000	0.000000	0.000000
75%	71.000000	3.760000	0.000000	1.000000
max	89.000000	5.000000	1.000000	1.000000
Mode	72.000000	0.690000	0.000000	0.000000

	Diagnosis
count	212691.000000
mean	0.232708
std	0.422559
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000
Mode	0.000000

```
In [95]: # Define selected variables for outlier detection (excluding Radiation_Exposure)
selected_vars_for_outliers = ["Age", "Nodule_Size", "Family_History"]

# Define function for IQR-based outlier detection
def detect_outliers_iqr(data, column):
    Q1 = data[column].quantile(0.25)
    Q3 = data[column].quantile(0.75)
    IQR = Q3 - Q1 # Compute IQR
    lower_bound = Q1 - 1.5 * IQR # Lower bound for outliers
    upper_bound = Q3 + 1.5 * IQR # Upper bound for outliers
    outliers = data[(data[column] < lower_bound) | (data[column] > upper_bound)]
    return Q1, Q3, IQR, lower_bound, upper_bound, len(outliers)

# Compute IQR and outliers for selected variables
iqr_results = {}
for col in selected_vars_for_outliers:
    Q1, Q3, IQR, lower_bound, upper_bound, num_outliers = detect_outliers_iqr(data, col)
    iqr_results[col] = {
        "Q1": Q1, "Q3": Q3, "IQR": IQR,
        "Lower Bound": lower_bound, "Upper Bound": upper_bound,
        "Num Outliers": num_outliers
    }

# Convert results to a DataFrame for better visualization
iqr_outlier_df = pd.DataFrame(iqr_results).T

# Print IQR-based Outlier Analysis
print("\n### IQR-Based Outlier Analysis (Excluding Radiation Exposure and Diagnosis) ###")
print(iqr_outlier_df)
```

```

### IQR-Based Outlier Analysis (Excluding Radiation Exposure and Diagnosis)
###

```

	Q1	Q3	IQR	Lower Bound	Upper Bound	Num Outliers
Age	33.00	71.00	38.00	-24.000	128.000	0.0
Nodule_Size	1.25	3.76	2.51	-2.515	7.525	0.0
Family_History	0.00	1.00	1.00	-1.500	2.500	0.0

```

In [97]: # Generate histograms

plt.figure(figsize=(15, 10))
colors = ["blue", "red", "green", "purple"]

for i, var in enumerate(selected_vars):
    plt.subplot(2, 3, i + 1)

    if var in ["Radiation_Exposure", "Family_History"]:
        sns.countplot(x=df_thyroid_encoded[var], palette=colors[i % len(colors)])
        plt.xlabel(f"{var} (0=No, 1=Yes)")
        plt.ylabel("Count")

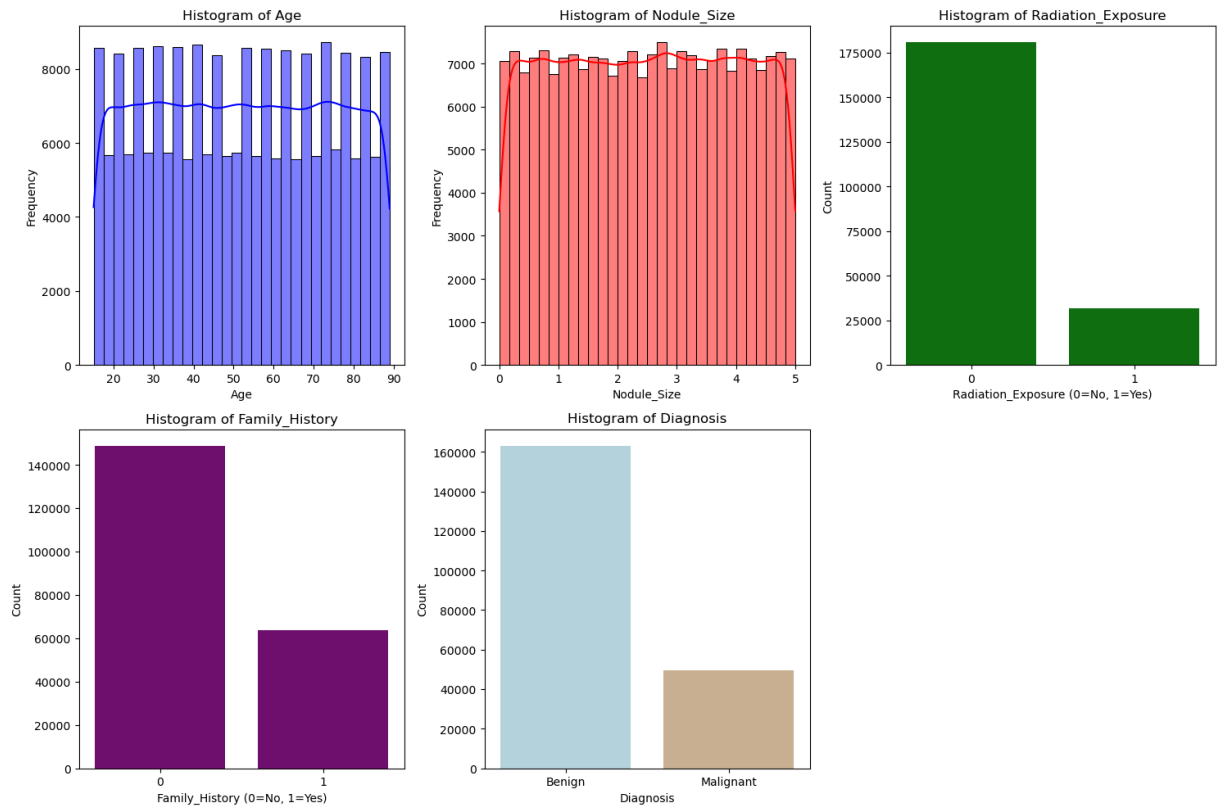
    elif var == "Diagnosis":
        sns.countplot(x=df_thyroid_encoded[var], palette=["lightblue", "tan"])
        plt.xlabel("Diagnosis")
        plt.ylabel("Count")
        plt.xticks(ticks=[0, 1], labels=["Benign", "Malignant"])

    else: # Continuous variables (Age, Nodule Size)
        sns.histplot(df_thyroid_encoded[var], bins=30, kde=True, color=colors[i % len(colors)])
        plt.xlabel(var)
        plt.ylabel("Frequency")


    plt.title(f"Histogram of {var}")

plt.tight_layout()
plt.show()

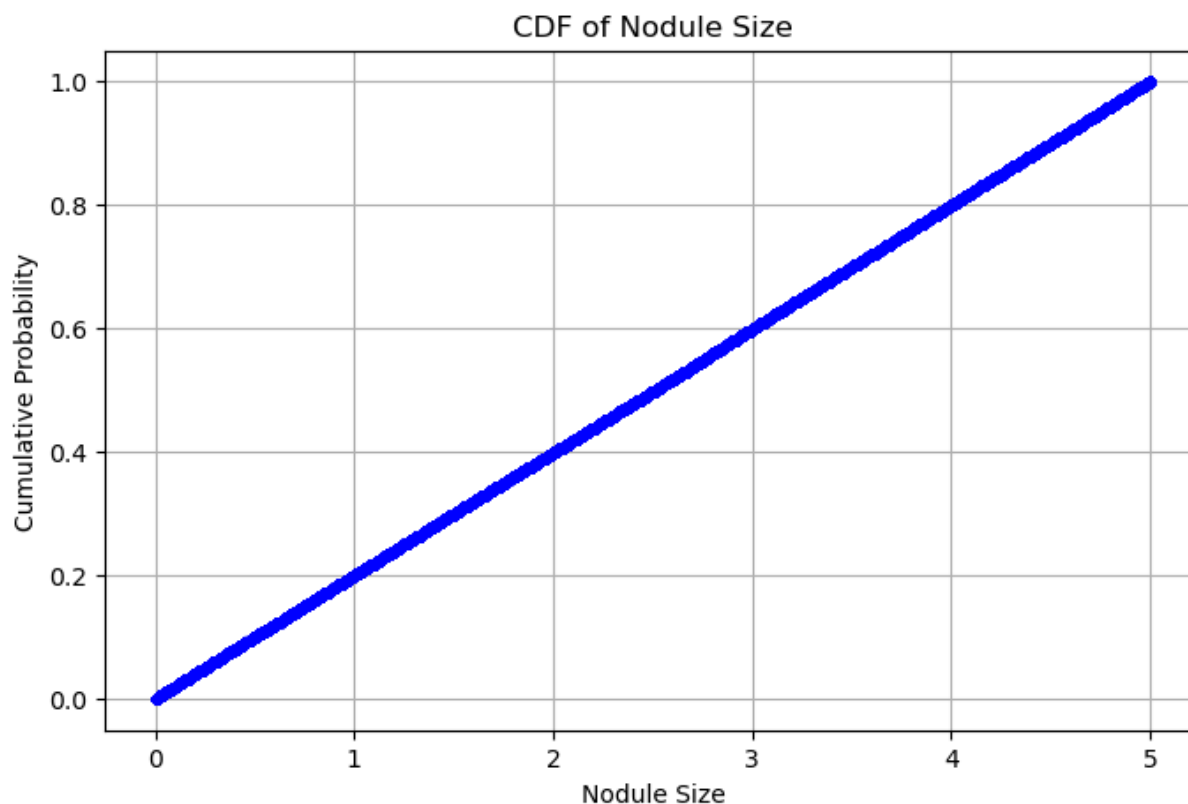
```



```
In [99]: # PMF Analysis for Nodule Size
df_thyroid_encoded["Nodule_Size_Binned"] = pd.cut(df_thyroid_encoded["Nodule_Size"],
pmf_benign = df_thyroid_encoded[df_thyroid_encoded["Diagnosis"] == 0]["Nodule_Size_Binned"],
pmf_malignant = df_thyroid_encoded[df_thyroid_encoded["Diagnosis"] == 1]["Nodule_Size_Binned"],
plt.figure(figsize=(8, 5))
plt.bar(pmf_benign.index, pmf_benign, alpha=0.6, label="Benign", color="blue")
plt.bar(pmf_malignant.index, pmf_malignant, alpha=0.6, label="Malignant", color="tan")
plt.xlabel("Nodule Size Groups")
plt.ylabel("Probability Mass Function (PMF)")
plt.title("PMF Comparison of Nodule Size for Benign vs. Malignant Cases")
plt.legend()
plt.show()
```



```
In [101... # CDF Analysis
sorted_nodule_size = np.sort(df_thyroid_encoded["Nodule_Size"])
cdf_nodule_size = np.arange(1, len(sorted_nodule_size) + 1) / len(sorted_nodule_size)
plt.figure(figsize=(8, 5))
plt.plot(sorted_nodule_size, cdf_nodule_size, marker=".", linestyle="none",
plt.xlabel("Nodule Size")
plt.ylabel("Cumulative Probability")
plt.title("CDF of Nodule Size")
plt.grid(True)
plt.show()
```



```
In [103... # Analytical Distribution: Normal Fit to Age

# Fit a normal distribution to Age
mu, sigma = np.mean(df_thyroid_encoded["Age"]), np.std(df_thyroid_encoded["Age"])

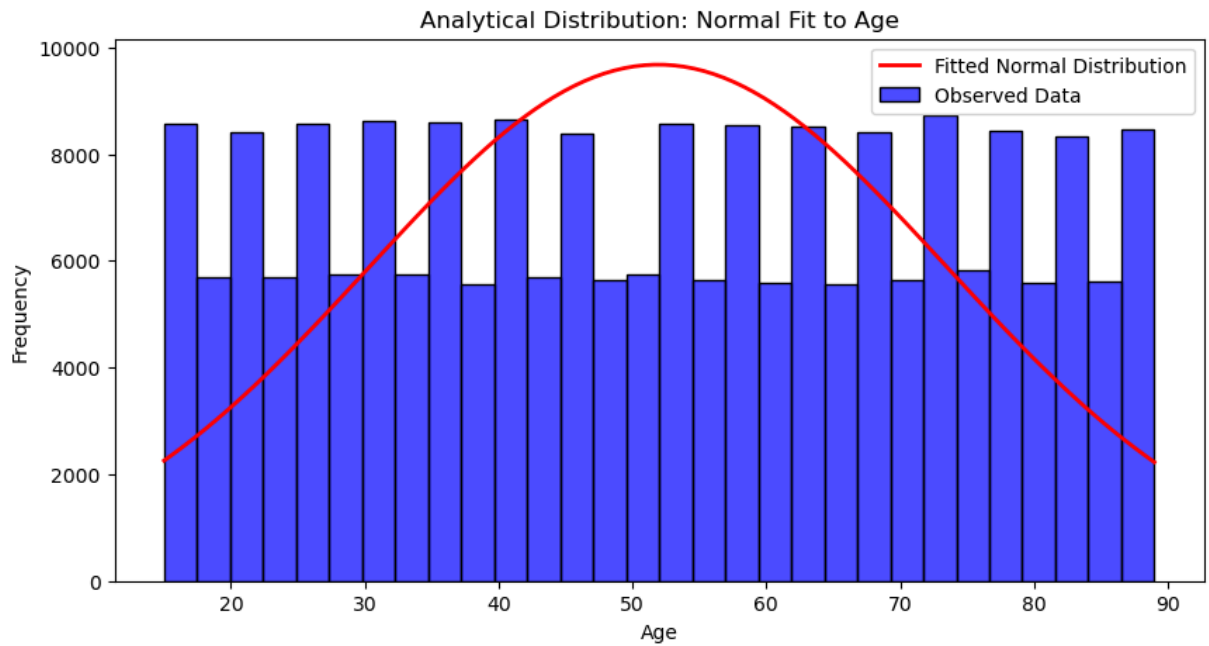
# Generate normal distribution curve
x = np.linspace(min(df_thyroid_encoded["Age"]), max(df_thyroid_encoded["Age"]), 100)
pdf = stats.norm.pdf(x, mu, sigma) * len(df_thyroid_encoded["Age"]) * np.diff(x)

# Plot histogram of observed Age data
plt.figure(figsize=(10, 5))
sns.histplot(df_thyroid_encoded["Age"], bins=30, kde=False, color="blue", alpha=0.5)

# Overlay the fitted normal distribution curve
plt.plot(x, pdf, color="red", linewidth=2, label="Fitted Normal Distribution")

# Titles and labels
plt.title("Analytical Distribution: Normal Fit to Age")
plt.xlabel("Age")
plt.ylabel("Frequency")
plt.legend()

# Show plot
plt.show()
```

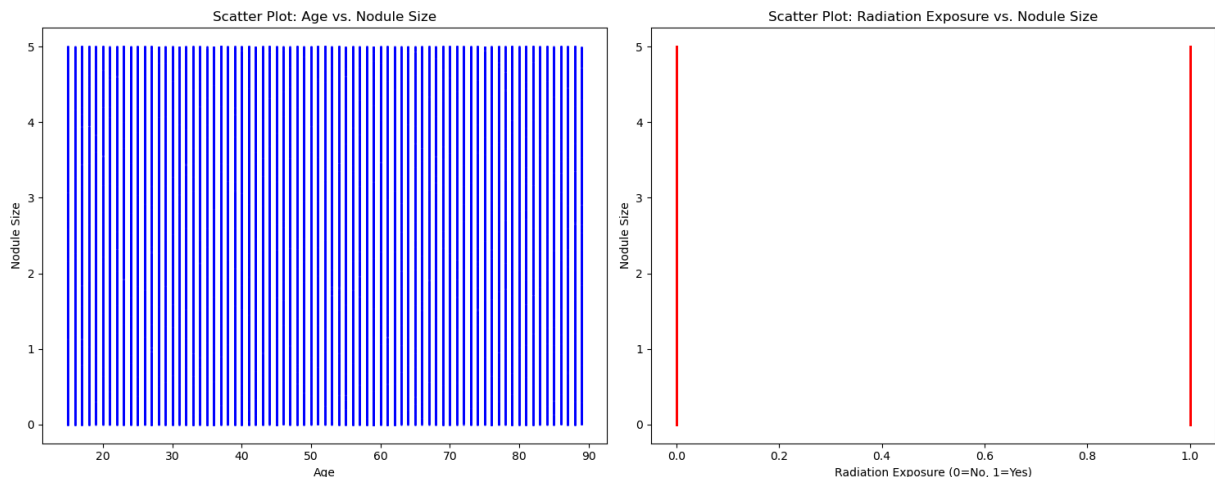



```
In [105... # scatter plots
# Define figure size
fig, axes = plt.subplots(1, 2, figsize=(15, 6))

# Scatter plot for Age vs. Nodule Size
axes[0].scatter(df_thyroid_encoded["Age"], df_thyroid_encoded["Nodule_Size"])
axes[0].set_xlabel("Age")
axes[0].set_ylabel("Nodule Size")
axes[0].set_title("Scatter Plot: Age vs. Nodule Size")

# Scatter plot for Radiation Exposure vs. Nodule Size
axes[1].scatter(df_thyroid_encoded["Radiation_Exposure"], df_thyroid_encoded["Nodule_Size"])
axes[1].set_xlabel("Radiation Exposure (0=No, 1=Yes)")
axes[1].set_ylabel("Nodule Size")
axes[1].set_title("Scatter Plot: Radiation Exposure vs. Nodule Size")

# Adjust layout and display
plt.tight_layout()
plt.show()
```



```
In [107... # Compute Pearson Correlation Coefficient
selected_vars = ["Age", "Nodule_Size", "Radiation_Exposure", "Family_History"]
pearson_corr_all = df_thyroid_encoded[selected_vars].corr(method="pearson")

# Print Pearson Correlation Coefficient Matrix
print("\n### Pearson Correlation Coefficient Matrix (All Selected Variables)
print(pearson_corr_all)
```

```
### Pearson Correlation Coefficient Matrix (All Selected Variables) ###
      Age  Nodule_Size  Radiation_Exposure  Family_History
y \
Age      1.000000    -0.001489           0.004007           0.00333
7
Nodule_Size -0.001489     1.000000           0.000118          -0.00287
7
Radiation_Exposure 0.004007     0.000118           1.000000          -0.00079
5
Family_History    0.003337    -0.002877          -0.000795           1.00000
0
Diagnosis      0.000115    -0.002658           0.089043           0.14092
1

      Diagnosis
Age      0.000115
Nodule_Size -0.002658
Radiation_Exposure 0.089043
Family_History 0.140921
Diagnosis      1.000000
```

```
In [109... # Hypothesis Analysis Chi-Square Test for Radiation Exposure vs. Diagnosis
contingency_table = pd.crosstab(df_thyroid_encoded["Radiation_Exposure"], df_thyroid_encoded["Diagnosis"])
chi2, p, _, _ = stats.chi2_contingency(contingency_table)
print("\n --- Chi-Square Test Results --- ")
print(f"Chi-Square Statistic: {chi2}, p-value: {p}")
```

```
--- Chi-Square Test Results ---
Chi-Square Statistic: 1685.7536394973827, p-value: 0.0
```

```
In [111... # Logistic Regression Analysis
X = sm.add_constant(df_thyroid_encoded[["Age", "Nodule_Size", "Radiation_Exposure"]])
y = df_thyroid_encoded["Diagnosis"]
logit_model = sm.Logit(y, X).fit()
print("\n ----Logistic Regression Summary ----")
print(logit_model.summary())
```

Optimization terminated successfully.
Current function value: 0.529167
Iterations 5

----Logistic Regression Summary ----					
Logit Regression Results					
=====					
==					
Dep. Variable:	Diagnosis	No. Observations:	2126		
91					
Model:	Logit	Df Residuals:	2126		
86					
Method:	MLE	Df Model:			
4					
Date:	Sun, 02 Mar 2025	Pseudo R-squ.:	0.024		
63					
Time:	22:26:21	Log-Likelihood:	-1.1255e+		
05					
converged:	True	LL-Null:	-1.1539e+		
05					
Covariance Type:	nonrobust	LLR p-value:	0.0		
00					
=====					
=====					
	coef	std err	z	P> z	[0.025
0.975]					

const	-1.5072	0.017	-88.944	0.000	-1.540
-1.474					
Age	-8.448e-05	0.000	-0.351	0.726	-0.001
0.000					
Nodule_Size	-0.0038	0.004	-1.042	0.297	-0.011
0.003					
Radiation_Exposure	0.5535	0.013	41.194	0.000	0.527
0.580					
Family_History	0.6953	0.011	64.643	0.000	0.674
0.716					
=====					
=====					

In []: